

Kestrel AI



Low-latency inference infrastructure for real-time AI

Founder & CEO: Jana Kessler · Munich

PROBLEM

Real-time AI is bottlenecked by inference latency

- Teams wait 300ms+ per call on generic GPU serving
- Autoscaling is manual, costly, and brittle
- No tooling built for sub-50ms production inference

PRODUCT

A serving layer purpose-built for latency

- Drop-in runtime: 40ms p95 on standard hardware
- Automatic batching + cache for repeated inference
- Open-source core, enterprise control plane

TRACTION

Real revenue, growing

€40K ARR

- 4 enterprise pilots (2 converted to paid)
- Launched 8 months ago; 2,100 GitHub stars
- Net revenue retention 118% across pilots

TEAM

Deep technical founder

- Jana Kessler — ex-DeepMind research engineer
- 2 papers at NeurIPS on efficient inference
- Prior: shipped ML infra used by 40M+ users

MARKET

AI inference tooling is a \$14B market by 2028

- Every AI product needs low-latency serving
- Bottom-up: 60k AI teams x €20k ACV
- Wedge: latency-critical realtime use cases

ASK

Raising €1.5M pre-seed

- 18 months runway; 4 engineering hires
- Target: 25 paying enterprises, €500K ARR